

AI Should Not Be Called Human-Like: An Approach From Conceptual Engineering

August 2026

ECSP; AI Ethics

Prof. Christoph Kelp and Prof. Mona Simion

Camilo Salgar

The question whether AI poses an existential risk has been framed as a question whether AI will disempower humanity such that humans cannot determine their own future (Cappelen et al., 2025; Dung, 2025). However, this restricts the existential problem to a matter of power, disregarding other forms of existential risk. Another way to frame the existential problem, which I shall develop in this paper, is that the development of AI changes what it means to be human. If AI poses a threat to humanity's understanding of itself, it will also lead to an existential problem (I will refer to this problem as the self-understanding problem). To pursue such an angle, conceptual engineering provides important insights since it cares precisely how we *should* define concepts to make sure they align with what we hope to accomplish. While conceptual engineering often restricts itself to discussing criteria and accuracy for certain cases, it is also meant to address social consequences of concepts (Löhr, 2024). Approaching the self-understanding problem through conceptual engineering, this paper argues that establishing a relationship of *likeness* between AIs and humans (as it has demonstrably been done by several researchers) is both inaccurate and raises the self-understanding problem. Consciousness will be constantly referred to as one of the main examples of the discourse on AI and human *likeness*; yet, this argument is not restricted to consciousness, but applies to any feature that humans have that is not trivially argued to be present in AI. This paper will follow a clear structure: I will, first, establish precisely how the self-understanding problem is an existential risk; second, entertain the general argument that AI is human-like in some respects; third, break down why this is an inaccurate argument and how it raises the self-understanding problem; finally, address counter arguments defenders might pose.

Any phenomena that pose a significant threat to what humans understand themselves to be (the self-understanding problem) are existential threats. Dung (2025) and Cappelen et al. (2025) both frame the existential threat of AI as a disempowerment problem. While that is an important consideration, it fails to acknowledge that there are other ways in which AI may pose an existential problem. It is useful to establish that what is referred to as an existential threat to humanity is a threat that may put humanity, as we know it, at risk. Notice that through this view it is no longer necessary that this threat be physical disempowerment; any phenomenon that may significantly change humanity should be considered an existential risk. One way that humanity can significantly change, is if what it means to be human changes. Considering an extreme case, if what it means to be human goes from being free to being subjugated, an existential problem arises. What it means to be human, in this paper, is not a metaphysical question but a social one. If the social understanding of what it means to be human is significantly disrupted, an existential risk arises; if a phenomenon can enact such disruption, it poses an existential risk.

Whenever some new phenomenon takes place, researchers play an important role in scrutinizing how such phenomena fit into preexisting concepts and structures. In social media and informal conversation, anthropocentric concepts are constantly attributed to AI (Coghlan & Antonellos, 2026). Indisputably, AI has exhibited the sort of phenomena that present challenges to our preexisting concepts; for instance: consciousness, agency, cognition, among others (Hillenbrink & Löhr, 2025). The process of examining what concepts mean and what they should mean is known as conceptual engineering. Conceptual engineering, Löhr (2024) argues, is not solely about technically assessing the criteria for when a concept is valid or not, but also the social consequences of such criteria. Thus, in the case of the concept of human-like being

attached to AI, it is relevant to examine the validity of such a concept and whether employing it raises the self-understanding problem.

First, we must evaluate what *likeness* between humans and AI means technically; then, why such *likeness* is inaccurate and how it may have negative social consequences.

Since an AI-human *likeness* relationship is meant to add information about AI by mapping qualities from humans to AI, it is presupposed that we know what that feature consists of in humans. Saying that AI, for some feature(s), is human-like, has the goal of adding the concept human-like to some feature in AI. Yet, for human-like to carry meaningful content about some feature, there must exist a reference for what that feature is in humans. Saying that cognition in AI is human-like (Cappelen & Dever, 2025), is a meaningful statement only if there is a clear standard to which human cognition is being held to be. Hence, when humans are a reference for a certain feature, a standard for such feature is presupposed. As an analogy, saying that Bob's hair is like Josh's hair carries no valuable information if no one knows what Josh's hair looks like. In the case of hair, simply looking at Josh may suffice for a standard. But for a concept like human consciousness, the equivalent is a theory for human consciousness. Notice that it does not matter whether the claim is that an entity has a human-like feature or that it is possible for such an entity to have a human-like feature, both presuppose a standard for such feature in humans. Thus, using the human-like descriptor for a certain feature implies establishing a standard of what the human feature is.

Theories based on anthropocentric abstraction are used to argue that AI has human-like features; however, using anthropocentric abstraction as a standard for what a human feature is, is circular logic. As Cappelen and Dever (2021) point out, anthropocentric abstraction begins by identifying a feature of humans with certain particularities and then reaching an abstractive zone

that generalizes the feature by removing anthropocentric features. Using their language for reference, Cappelen and Dever outline the process of anthropocentric abstraction as follows, “We start with a pattern instantiated by human peculiarities. There is then a hierarchy of degrees of abstractions: what we are looking for is a degree of abstraction that preserves sufficiently many important features of the original phenomenon. It’s not too abstract and it’s not too focused on specific details. Call that ‘the abstractive sweet-spot’” (Cappelen & Dever, 2021, p. 87). Appropriate anthropocentric abstraction theories are ones that successfully lie in the abstractive sweet spot. Such theories are tested with counterexamples to make sure that it accurately tracks consciousness; for instance, a theory that would establish that a phone is conscious, must reconsider its abstraction. (Cappelen & Dever, 2021). In this process, the fundamental premise is that humans are conscious; if that were not the case, the anthropocentric abstraction would make no sense. A limitation for theories derived as such is that they hold no explanatory bearing on humans. The circularity goes as follows: (1) humans are conscious (2) anthropocentric abstraction can be established for consciousness by reducing human particularities (3) whichever entity meets the criteria of the anthropocentric abstraction can be said to be conscious (4) humans meet the criteria (5) humans are conscious. Thus, a theory based on anthropocentric abstraction cannot avoid circularity when used to ground the human feature it was based on.

The descriptor human-like cannot be validly used to describe features in non-human entities (or other AI features) when the justification is based on anthropocentric abstraction. As stated above, a theory based on anthropocentric abstraction cannot supply an answer to the question of what human consciousness is. Additionally, saying AI has some human-like features requires a standard for what that feature is in humans. To be exact, a standard is only appropriate if it accurately grounds a feature. For instance, in the case of saying that Bob’s hair is like Josh’s

hair, a standard for Josh's hair is only useful if it explains how Josh's hair looks. We can now string these claims together: (C1) an anthropocentric abstraction cannot non-circularly ground a standard for the feature as it is in humans; (C2) saying that AI has a human-like feature requires an accurate standard of what that human feature is. Combining C1 and C2, a *likeness* relationship between AI and humans cannot be justified through anthropocentric abstraction.

A *likeness* relationship between AI and humans is not only invalid, but it also raises the self-understanding problem since what may start as a one-way *likeness* relationship, becomes two-way as soon as a social effect is considered. As explained above, when a researcher says that AI is human-like regarding some feature, the researcher has the intention of mapping some understanding of humans onto AI. However, the opposite may also take place, such that an understanding of AI is mapped onto humans. Consider the following thought experiment:

Let Edward, Herbert, and Alice be three people. Suppose Edward knows both Herbert and Alice well, and Edward considers them to have similar personalities. One day, Edward is eating dinner with a group of random people who know Herbert, and Edward mentions Alice, but they have no idea who Alice is. To clarify, Edward says, "Alice's personality is like Herbert's". The group of random people, knowing that Edward knows both Alice and Herbert well enough, take Edward's word for it; thus, they effectively gain information about Alice using Herbert as reference. However, sometime later, the random people observe Alice being dishonest and mean and recall Edward saying that Alice's personality is like Herbert's. Considering that the random people are aware Edward knows both Herbert and Alice well, and they have no reason to distrust Edward, they now have reason to believe either that Herbert is not the respectful and honest person they believed Herbert to be or that Edward should not be trusted.

Edward represents the role of experts in the AI discourse, Herbert the concept of human, and Alice, AI; saying Alice's personality is like Herbert's is analogous to saying AI is human-like for some feature. The objective of this Edward-Herbert-Alice thought experiment is to demonstrate that a claim of the form X is like Y, in a social environment, may also transfer information from X to Y. If the trustworthiness of the interlocutor is difficult to doubt, claim X is like Y allows for information to be transferred from X to Y in a social environment.

For further clarity it is helpful to explain that this is a social argument, not an abstract one. Abstractly, Tversky (1977) made it clear that the proposition subject X is like referent Y is not the same as referent Y is like subject X, simply because the direction of subject-referent *likeness* matters. For instance, saying humans are like machines when it comes to following instructions is different from saying machines are like humans when it comes to following instructions; the former indicates that humans are good at following instructions, whereas the latter means that machines can decide not to follow instructions. However, these logical games presuppose perfect information for the subject and the referent; the social effect takes place when the people accepting that X is like Y have imperfect information. Returning to the Edward-Herbert-Alice thought experiment, Edward's standing as friend of both Alice and Herbert gives people reason to believe that Edward already knew the information they gained later. Thus, the process of gaining information that the expert is thought to already have, makes it so the X is like Y relation leads to a transfer of knowledge from X to Y.

The social effect of *likeness* relationships poses a threat to our understanding of what humans are. An expert is socially held as a trustworthy authority. So, when an expert in the philosophy of AI says that some feature of AI is human-like, it enters the social effect noted above. It turns out, the descriptor human-like is a device that changes the concept of human.

Consider the following thought experiment: John, someone who has never interacted with AI chatbots, hears an expert claim that AI cognition is human-like. Intrigued, John interacts with an AI chatbot like Claude (Anthropic, 2026), and John realizes how Claude breaks down its answers step-by-step and is extremely rational. Since John trusts the expert who claimed that AI cognition is human-like, it is reasonable for him to observe how Claude cognizes and adjust his own understanding of what human cognition is. John may now understand human cognition as extremely rational. At a large enough scale, this may significantly change the social understanding of what humans are. A *likeness* relationship tethers a social understanding of what humans are to observations of AI. If what it means to be human varies depending on features observed in AI, then the *likeness* relationship may raise the self-understanding problem.

Some researchers argue that the term like is useful since it acknowledges both consciousness in AI and humans while recognizing undoubtable particularities in each; however, this is a mistaken use of the term and falls into the problems described above. Theories of consciousness based on anthropocentric abstraction have established the possibility of AI being conscious (Chalmers, 2023). However, saying that consciousness in AI would be the same as consciousness in humans is not justified since there are irreconcilable differences between AI and humans (Krzanowski, 2026). Thus, it is assumed that saying AIs have human-like consciousness is a much more accurate term. The term like is meant to establish that according to some theory, AIs and humans may both be conscious, while recognizing that there are particularities in AIs and humans that cannot be completely neglected. In Krzanowski's own words: "we need to recognize that in constructing AI systems (because of the gap [between humans and AI]) we are not building synthetic human systems but rather synthetic human-like systems" (2026, p. 276). Krzanowski argues that the -like suffix, while vague, does the work of

acknowledging the gap between AIs and humans such that it is not completely neglected. Note that I am not arguing that it is problematic to state that AI has consciousness according to some theory; the problem arises when the statement is that AI has human-like consciousness according to a theory. However, the problem of the -like suffix is not that it recognizes the gap; rather, that it establishes a relationship across the gap such that observing AI may change our understanding of humans.

Taking both the validity of a concept and its social effect into account, the statement that AI has or may have human-like features based on anthropocentric abstraction is invalid and raises the problem of self-understanding. Even if a non-anthropocentric abstraction sustained AI and human *likeness*, implementing such a term should still be discussed since the self-understanding problem may still arise. Non-trivial AI features (e.g. consciousness, cognition, agency) should be discussed on their own terms, without a *likeness* to humans being established. It remains to be discussed whether the *likeness* relationship generates a significant enough change in what it means to be human for it to be an existential risk. Regardless, it is clear that we should be careful in how we approach AI philosophy, because depending on our approach we might be implicitly changing our understanding of humanity.

I shall now address two foreseeable objections. The objections will be voiced from the perspective of a critique, and the reply will be my own argumentative perspective.

Objection 1: It is unclear why allowing our understanding of AI to adjust our understanding of humanity is problematic; contrarily, AI may provide insights into what is going on in humans.

Reply 1: This objection makes the case that if two systems are alike, then examining one may prove insightful for the other. That is reasonable, yet it assumes that AIs and humans are alike. What is occurring in academic literature is that AI is providing insights on concepts that were

previously exclusively applied to humans. Making the distinction is critical: insights on consciousness may be discussed, but human consciousness should remain intact. In other words, the discourse on non-trivial AI features (e.g. consciousness) should separate consciousness from human consciousness, since AI being conscious does not necessarily mean it has human-like consciousness.

Objection 2: The Edward-Herbert-Alice thought experiment only establishes a mechanism through which information may flow contrary to the initial direction of a *likeness* relationship, it does not demonstrate that such a flow is problematic. Likewise, it is unclear why information flowing from AI to humans is problematic rather than insightful.

Reply 2: This is a strong objection, since my argument assumes that letting our understanding of humanity change based on AI is normatively undesirable. While a proper defense of this assumption requires an ethical investigation on whether it is risky for the social understanding of human features to be drawn towards AI features, this paper opts to step on the safer side. Intuitively, it is difficult to foresee what AI discoveries lie in the future (especially long-term), so keeping the understanding of humanity untethered from AI investigation will protect our self-understanding. This paper should serve as a warning to make sure researchers are aware of the potential repercussions of saying AI systems are human-like.

References

- Anthropic. (2026). Claude [Large language model]. <https://claude.ai/new>
- Cappelen, H., & Dever, J. (2021). *Making AI intelligible: Philosophical foundations*. Oxford University Press. <https://doi.org/10.1093/oso/9780192894724.001.0001>

- Cappelen, H., & Dever, J. (2025). Going whole hog: A philosophical defense of AI cognition. arXiv. <https://doi.org/10.48550/arXiv.2504.13988>
- Cappelen, H., Goldstein, S., & Hawthorne, J. (2025). AI survival stories: A taxonomic analysis of AI existential risk. *Philosophy of AI, 1*, 1–19. <https://doi.org/10.18716/ojs/phai/2025.2801>
- Chalmers, D. J. (2023). Could a large language model be conscious? arXiv. <https://doi.org/10.48550/arXiv.2303.07103>
- Coghlan, S., & Antonellos, M. (2026). Clarifying anthropomorphism for robots and AI: A critical review. *International Journal of Social Robotics*. Advance online publication. <https://doi.org/10.1007/s12369-026-01419-6>
- Dung, L. (2025). The argument for near-term human disempowerment through AI. *AI & Society*, 40(3), 1195–1208. <https://doi.org/10.1007/s00146-024-01930-2>
- Hillenbrink, R., & Löhr, G. (2025). Do new technologies change our moral norms? *Think*, 24(71), 55–59. <https://doi.org/10.1017/S1477175625100729>
- Krzanowski, R. (2026). Human nature and artificial intelligence: Sizing the gap. In V. C. Müller, L. Dung, G. Löhr, & A. Rumana (Eds.), *Philosophy of artificial intelligence* (Synthese Library, Vol. 533, pp. 275–287). Springer. https://doi.org/10.1007/978-3-032-10073-3_16
- Löhr, G. (2024). If conceptual engineering is a new method in the ethics of AI, what method is it exactly? *AI and Ethics*, 4(2), 575–585. <https://doi.org/10.1007/s43681-023-00295-4>
- Tversky, A. (1977). Features of similarity. *Psychological Review*, 84(4), 327–352. <https://doi.org/10.1037/0033-295X.84.4.327>